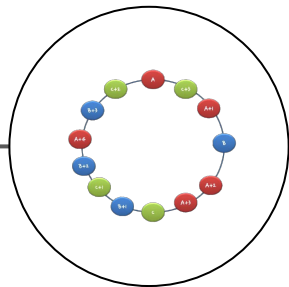


Lim, Hyeongseop

Backend Developer

Development Projects Overview

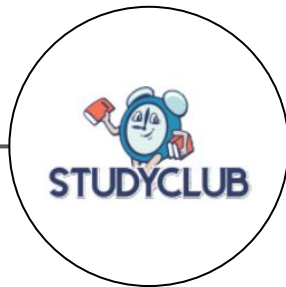


Consistent Hashing

Backend

Backend development project focusing on shard management and data redistribution systems.

2024.8 ~2024.9



STUDYCLUB

Backend

Web App for Study Recruitment and Activities: Unified platform for study groups.

2024.6 ~2024.7

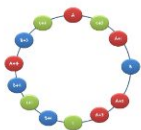


Doggiverse

Backend

Web App for Pet Owners: Content-based community service for pet owners and their pets.

2024.5 ~2024.6



Shard Management and Efficient Data Redistribution System

Development Period 2024.8 ~2024.9

Platform Web

Team Size 1

Development Environment

Languages Java(JDK 17), HTML/CSS, JavaScript

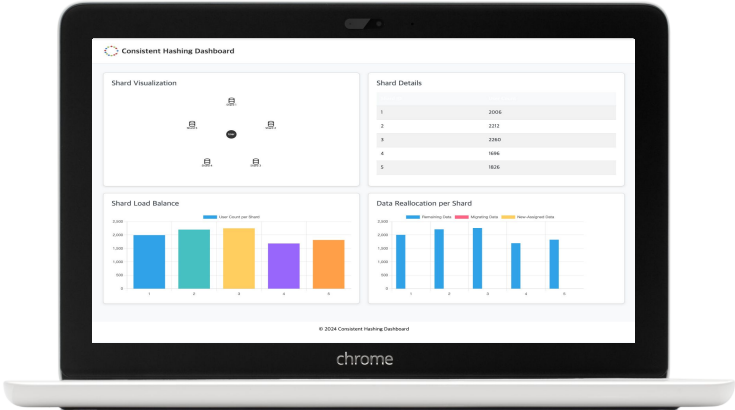
Server Spring Boot embedded Tomcat server

Frameworks Spring Boot, JPA, Spring WebSocket

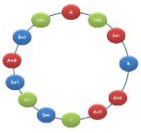
Database Oracle DB

IDE IntelliJ, SQL Developer, Postman

API, Libraries WebSocket API, Thymeleaf, Guava

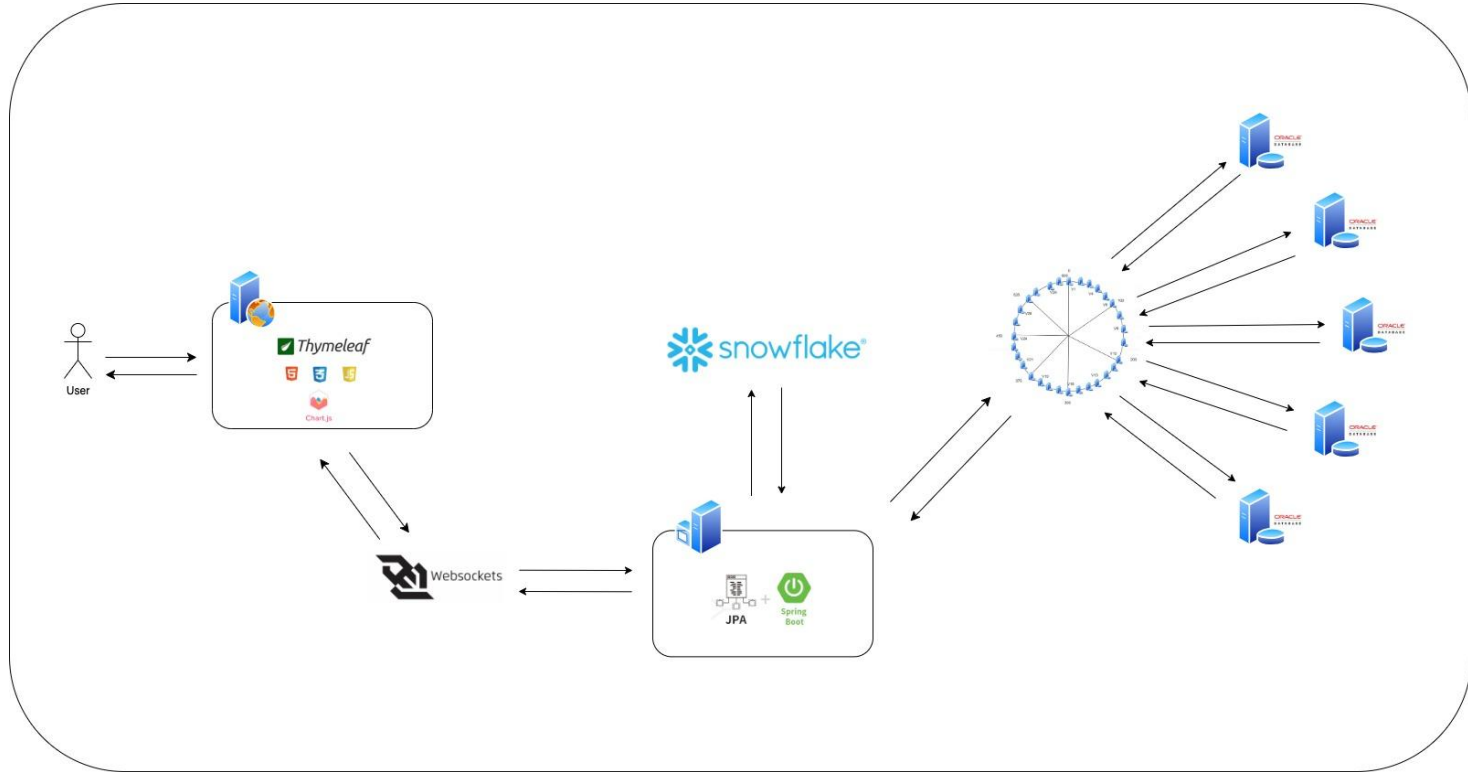


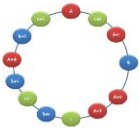
 **GitHub** <https://github.com/h3lim/Consistent-hashing>



Project Planning and Direction (1/4)

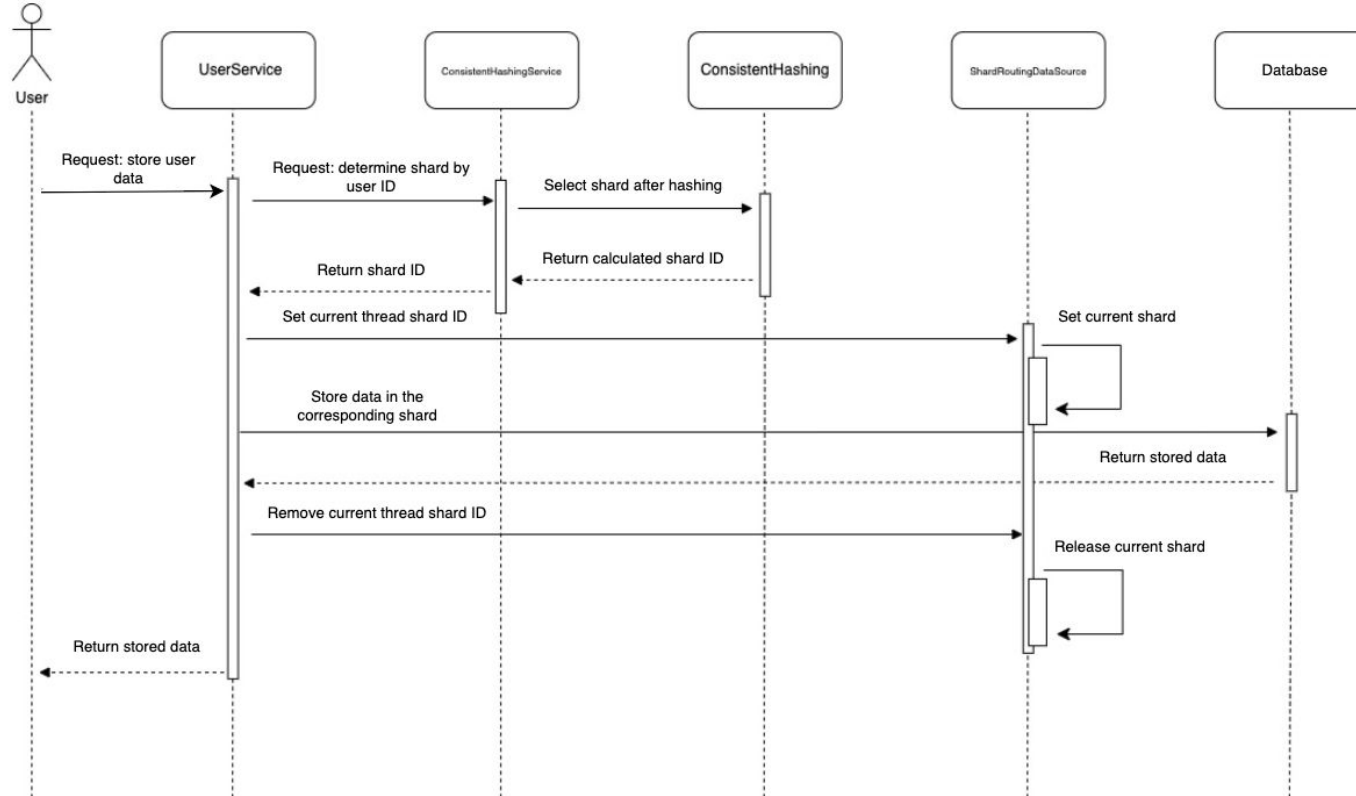
System Architecture

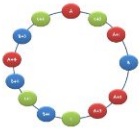




Project Planning and Direction (2/4)

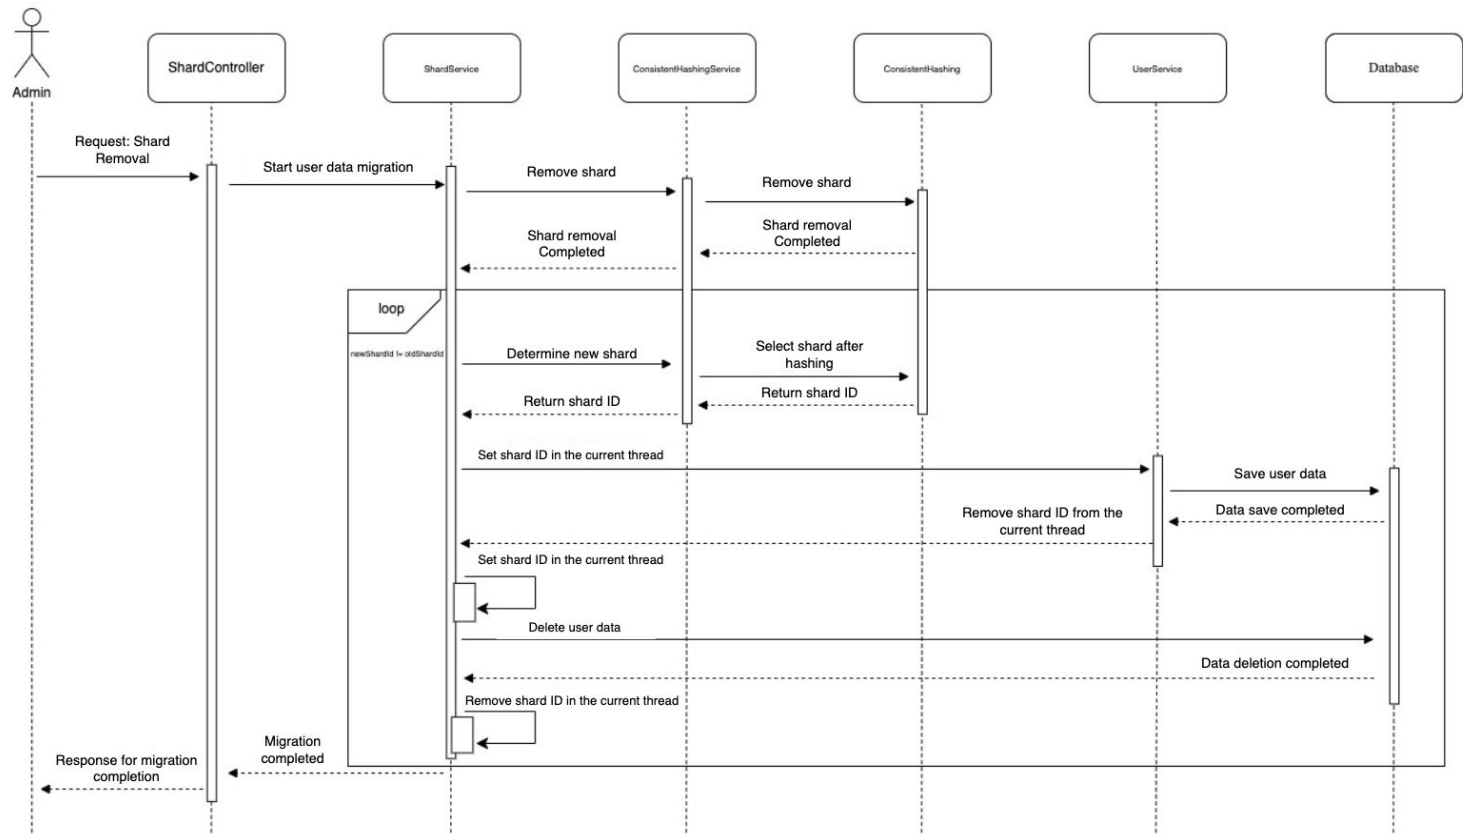
Sequence Diagram – Sharding Business Logic





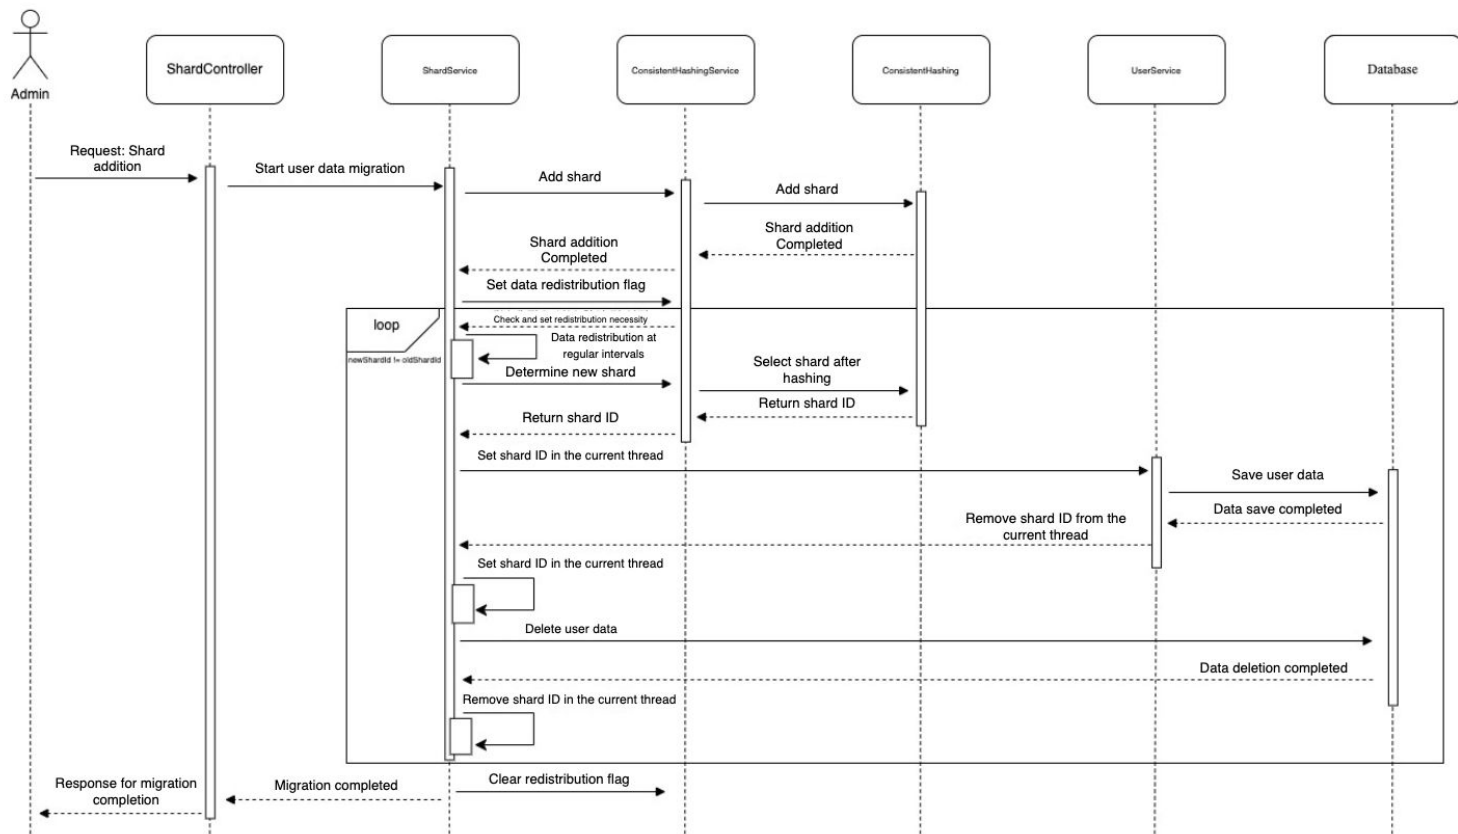
Project Planning and Direction (3/4)

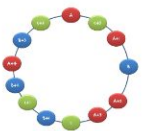
Sequence Diagram - Shard Removal and Data Redistribution





Sequence Diagram - Shard Addition and Data Redistribution





Key Features (1/3)

Core Sharding Logic

- Data is distributed using consistent hashing and stored across multiple database shards. This method allows for handling shard additions and removals with minimal data movement, avoiding the need to redistribute all data. This enables dynamic shard management while ensuring system scalability and stability.

Implementation Details

- ConsistentHashing Class:
 - The ConsistentHashing class manages the hash ring and calculates hashes using the given number of nodes (shards) and virtual nodes. It maps user IDs or data keys onto the hash ring using the murmur3_32 hashing algorithm.
 - addNode(int shard): A method that adds a shard to the hash ring based on virtual nodes.
 - removeNode(int shard): A method that removes a shard from the hash ring by deleting the virtual nodes associated with the removed shard.
 - getNode(long userId): A method that maps a given user ID or data key onto the hash ring and returns the corresponding shard.
- ConsistentHashingService Class:::
 - This service class determines which shard a user's data belongs to when storing or retrieving data from actual shards.
 - getShardForUser(Long userId): Takes a user ID as input, determines the shard the user belongs to, and manages data storage or retrieval in that shard.
 - addShard(int shardId): Adds a new shard to the hash ring and updates the shard distribution.



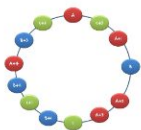
Key Features (2/3)

Unique Key Generation

- To generate unique identifiers for each user's ID or data record, a custom ID generator was implemented. This generator was inspired by the Snowflake algorithm and designed to create conflict-free identifiers in distributed systems by combining a unique node ID, a timestamp, and a sequence value.

Implementation Details

- CustomIdGenerator Class:
 - The node ID is a unique value assigned to each server or shard, ensuring that all IDs generated on that node are unique.
 - The timestamp is based on the current time, allowing the generation of multiple unique IDs even within the same time frame.
 - The sequence value prevents duplicate ID generation for the same timestamp.
- IDs are created by combining the millisecond-based timestamp, node ID, and sequence value, allowing the generation of millions of unique IDs in a short period.



Key Features (3/3)

Data Movement (Shard Migration)

- When shards are added or removed, data is redistributed to new shards. This redistribution is done using consistent hashing, which minimizes data movement and ensures asynchronous processing to prevent service interruptions. Additionally, real-time updates on data movement can be monitored through a dashboard using WebSocket.

Implementation Details

- **ShardService Class:**
 - `migrateUsersFromShard(int removedShardId)`: A method that migrates data from a removed shard to newly allocated shards.
 - `migrateUser(User user, int oldShardId)`: The logic for moving an individual user's data. When a user's shard is removed, this method locates a new shard and saves the user's data there.
- **WebSocket Integration:**
 - `sendShardUpdates()`: Periodically sends real-time updates on data distribution and movement via WebSocket, allowing administrators to monitor the dashboard.



Static Database Management Issues

During testing, repeated difficulties were encountered in registering and configuring data for each shard. Adding or removing shards required manual registration, which led to inconsistencies and prolonged test preparation times.

Symptoms

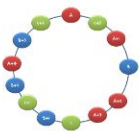
- Data registration was required for each database (shard) during every test.
- The more databases (shards) there were, the more time-consuming the manual setup became, leading to frequent errors during test environment setup.
- Manual configuration changes were needed each time a new shard was added or removed.

Cause Analysis

- Database settings were hard-coded, necessitating code modifications whenever a shard was changed, which was inefficient.
- The application lacked a feature for automatic shard registration or removal at runtime, making manual adjustments necessary for test environments.

Solutions

- Modified the system to use Spring's `@ConfigurationProperties` to dynamically load shard information from a configuration file, allowing automatic registration and management of shards as defined in the configuration during application startup.
- Used the `ShardDataSourceProperties` class to dynamically load and manage all shards.



Data Consistency Issues (During Shard Addition/Removal)

Issues were encountered when user data did not migrate correctly or data was lost during shard addition or removal, especially when handling large data volumes. This led to inconsistent data movement.

Symptoms

- Data disappearing instead of being transferred when a shard was removed.
- Existing data not being redistributed to new shards after a shard was added.
- Data inconsistencies, such as duplicated data across shards or data remaining unassigned.

Cause Analysis

- Inadequate transaction management during data movement, resulting in incomplete processing or failures that compromised data consistency.
- Attempting to move large amounts of data in real-time caused system load, leading to incomplete data migration.

Solutions

- Used transactions to ensure that data migration was only marked as successful when fully completed on a shard, with a rollback mechanism for failed transactions to maintain data consistency.
- Designed the migration process to move user data sequentially instead of all at once, reducing system load and ensuring stable data migration.



Web App for Study Recruitment and Activities

Development Period 2024.7 ~2024.8

Platform Web

Team Size 4

Responsibilities Implemented 1:1 video interviews using WebRTC, developed the quiz feature, and contributed to project planning and direction setting.

Development Environment

Languages Java(JDK 17), HTML/CSS, JavaScript

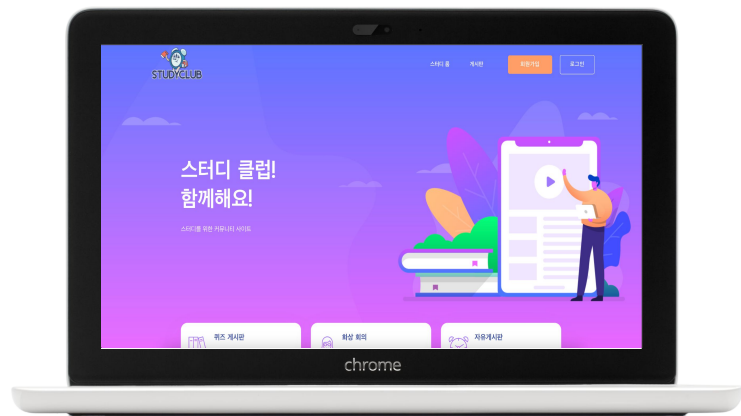
Server Apache Tomcat 9.0, WebSocket, WebRTC

Frameworks Spring MVC, Spring Boot, Spring Security, MyBatis 3.4.6

Database Oracle DB

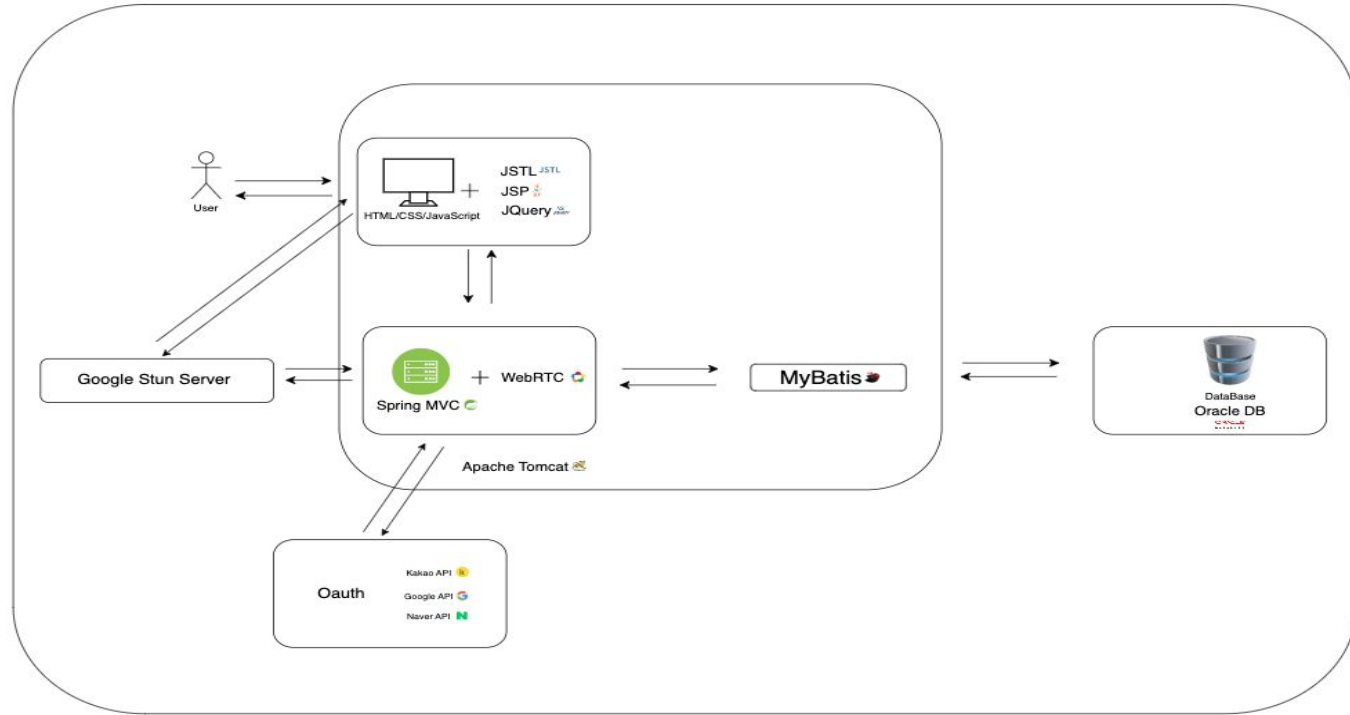
IDE IntelliJ, SQL Developer

API, Libraries KAKAO/NAVER/GOOGLE Login API, JQuery, WebRTC API, WebSocketAPI, Log4j, SLF4J, Validation

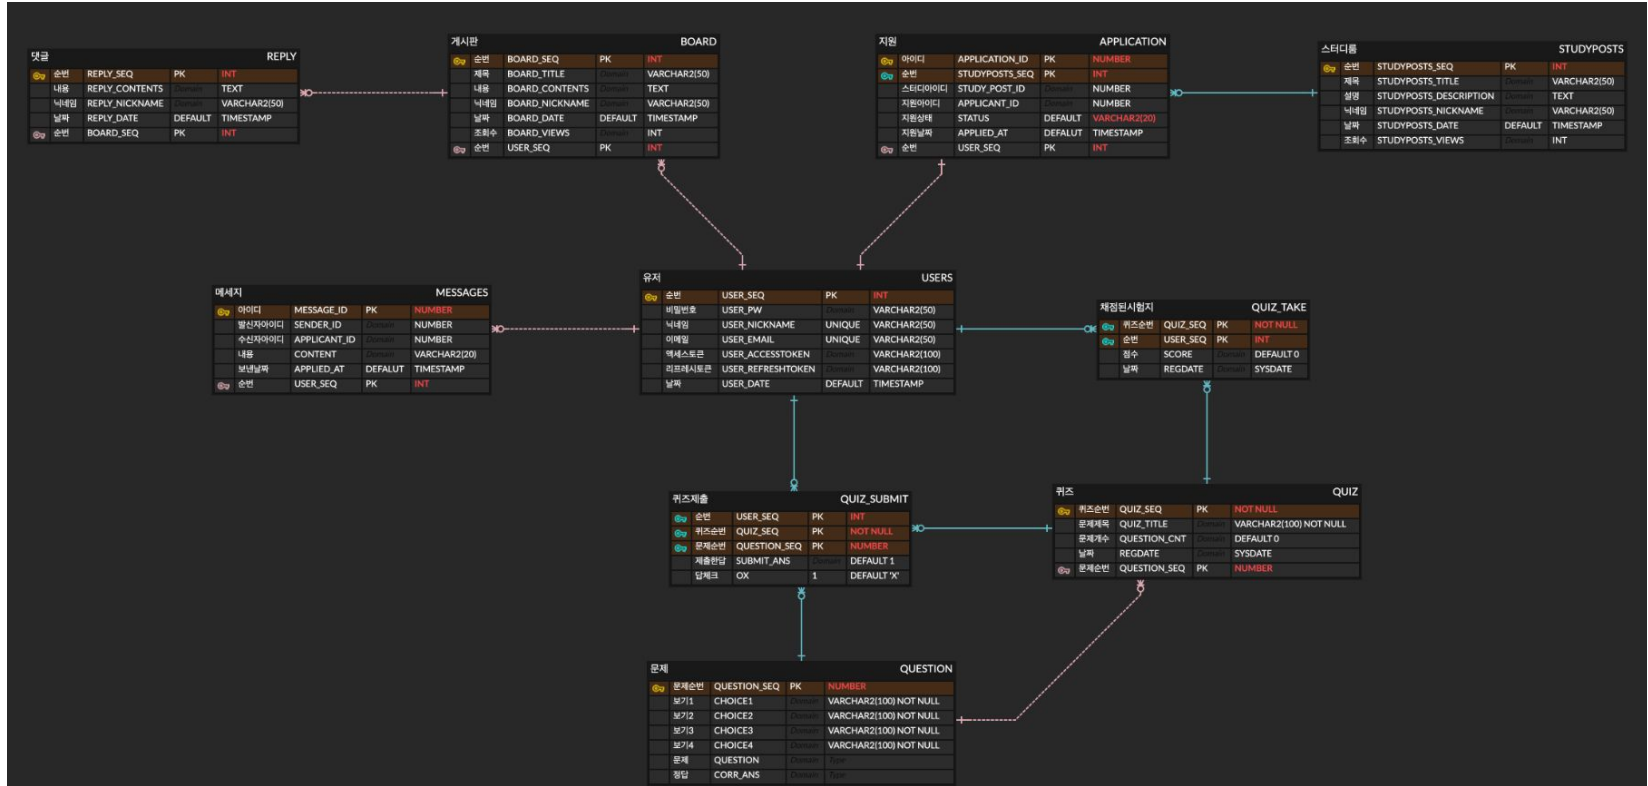


GitHub <https://github.com/h3lim/StudyClub>

System Architecture



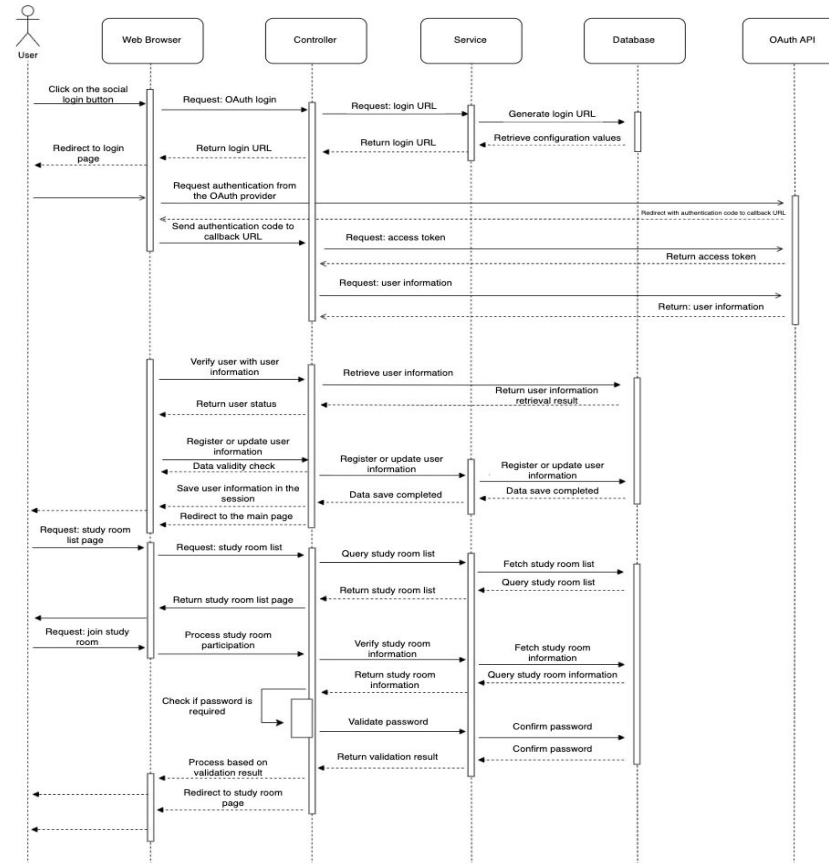
ERD Design



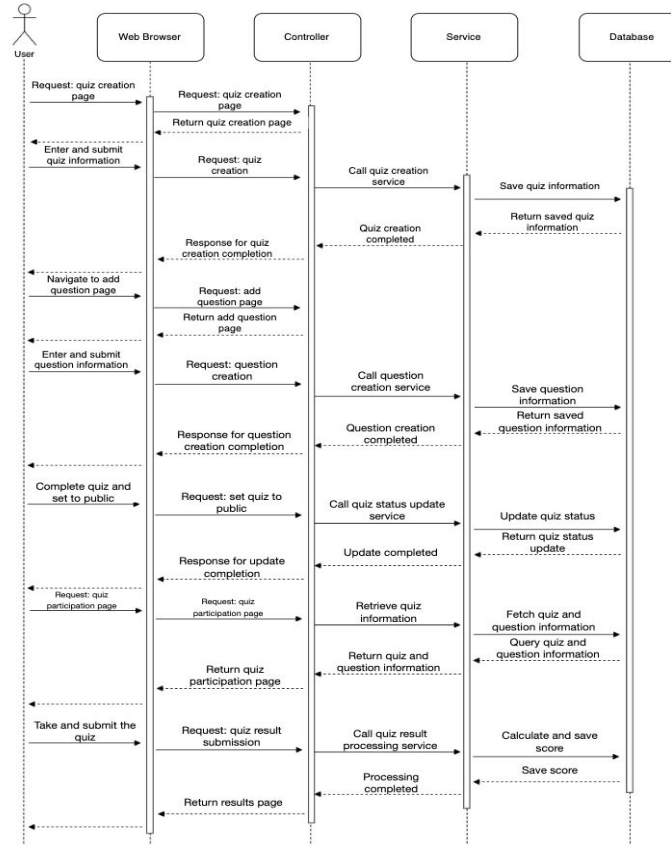


Project Planning and Direction (3/4)

Sequence Diagram – OAuth Login and Study Room Participation



Sequence Diagram – Quiz Creation and Participation





Real-Time Video Conference Functionality Using WebRTC

- In the STUDYCLUB project, WebRTC technology was used to enable real-time video meetings within study rooms, allowing remote users to communicate and collaborate as if they were in the same physical space.

Implementation Details

- Signaling Server Implementation: CamController was implemented to handle the signaling process necessary for establishing WebRTC connections.
- Use of STUN Server: A public STUN server was utilized to support connections between clients behind NAT firewalls.
- Communication via WebSocket: The WebSocket protocol was used for real-time message exchange between clients and the server.
- Client-Side Implementation:
 - Used JavaScript's RTCPeerConnection API to establish peer-to-peer connections.
 - Captured and transmitted user media streams (audio and video) to the counterpart.
 - Communicated with the server for signaling data exchange via WebSocket.



Quiz Creation and Participation Feature

- STUDYCLUB provides a feature for users to create quizzes and allows other users to participate in them to review study content, promoting knowledge sharing and learning motivation among users.

Implementation Details

- Quiz Creation Process:
 - Users input quiz titles and descriptions to create quizzes.
 - Multiple questions can be added to the created quiz.
 - Each question includes information such as the question text, choices, and the correct answer.
- Quiz Participation Process:
 - Other users can select and attempt public quizzes.
 - After answering and submitting the quiz, the score is calculated and displayed.
- Database Design:
 - Quiz Table: Stores basic information about quizzes.
 - Question Table: Stores questions related to each quiz.
 - UserGrade Table: Stores user quiz participation results.



Challenges in ERD Design for Quiz Feature

Significant time was spent designing the ERD (Entity-Relationship Diagram) for the quiz feature.

Challenges

- **Relationship Between Quizzes and Questions:** Each quiz contains multiple questions, and each question belongs to a specific quiz. How should this be represented efficiently?
- **Storing User Participation Results:** How to store and manage user participation data (scores, submitted answers, etc.) effectively?

Solutions

- **Quiz-Question Relationship:** Created Quiz and Question tables, with quiz_id in the Question table as a foreign key referencing quiz_id in the Quiz table to establish a clear 1:N relationship.
- **Storing User Participation Results:** Created a UserGrade table to store user quiz participation results, including columns like userNickname, quizSeq, score, and submitDate, allowing tracking of which user took which quiz and the score received.

Results

- A well-defined database structure facilitated efficient data processing for the quiz feature.
- Systematic management of user participation results enabled learning performance tracking and feedback.



Problem with Missing await in Asynchronous WebRTC Operations

Issues arose when the `send()` function was called without `await`, causing asynchronous task execution to become disordered. For example, `send()` was executed before `setLocalDescription` completed in the `createOffer()` function, leading to the transmission of incomplete signals.

Symptoms

- Offer and answer transmissions failed to function properly, causing unstable or failed connections.
- Audio or video streams were not sent, or real-time communication delays occurred.

Cause Analysis

- Calling `send()` without `await` led to signal transmission before `setLocalDescription` completed, resulting in unexpected signals on the receiver's side.
- Signal transmission timing became disordered, leading to unstable connections and missing signals.

Solutions

- Added `await` to ensure `send()` was executed only after the completion of asynchronous tasks, maintaining proper order.
- Adjusted `createOffer()` to send signals with `await send()` after `setLocalDescription` completed, ensuring correct signals were received.



Video Transmission Issues During 1:1 Interviews – HTTPS and SSL Certificate

Problems occurred with the WebRTC-based video call functionality during 1:1 interviews in STUDYCLUB, where the counterpart's video was not transmitted or frequently froze.

Symptoms

- Counterpart's video was not displayed, and WebRTC connections were blocked in the browser.
- WebRTC failure logs appeared in the console.

Cause Analysis

- WebRTC functions only over secure (HTTPS) connections. Without an SSL certificate or with a misconfigured certificate, WebRTC communication fails. Modern browsers block WebRTC over HTTP, causing video transmission failures.

Solutions

- Obtained and registered an SSL certificate in Tomcat and configured server.xml to enable HTTPS.
- Set up automatic redirection from HTTP to HTTPS.
- Configured the browser to access the application over HTTPS, ensuring secure WebRTC connections.

Video Transmission Issues During 1:1 Interviews – Port Forwarding and Firewall Problem

Problems were encountered with WebRTC video calls during 1:1 interviews in STUDYCLUB, where the video transmission was interrupted or frequently stopped.

Symptoms

- Counterpart's video was not transmitted or was interrupted intermittently.
- P2P connections were blocked due to certain ports being restricted by firewalls or routers.

Cause Analysis

- WebRTC communication relies on specific ports. Without proper port forwarding or if firewall rules block these ports, P2P connections fail, and video transmission is interrupted.
- Issues occurred when port forwarding on Tomcat server and Oracle DB was not properly configured or relevant ports (HTTP: 8087, HTTPS: 8447, Oracle DB: 1521) were blocked by firewalls.

Solutions

- Allowed ports (8447, 8087, 1521) necessary for WebRTC communication in firewall and router settings.
- Verified firewall settings to ensure port allowances and added rules to open these ports.
- Added specific ports to inbound firewall rules and applied them to TCP.

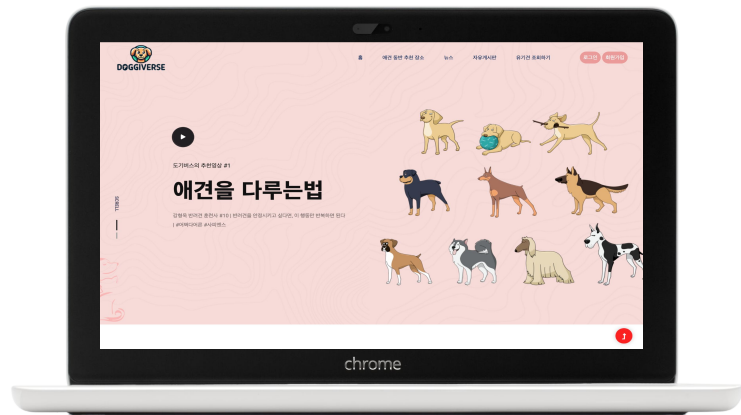


Web App for Pet Owners: Content-based community service

Development Period	2024.5 ~2024.6
Platform	Web
Team Size	3
Responsibilities	Implemented user registration and My Page features, , and contributed to project planning and direction setting

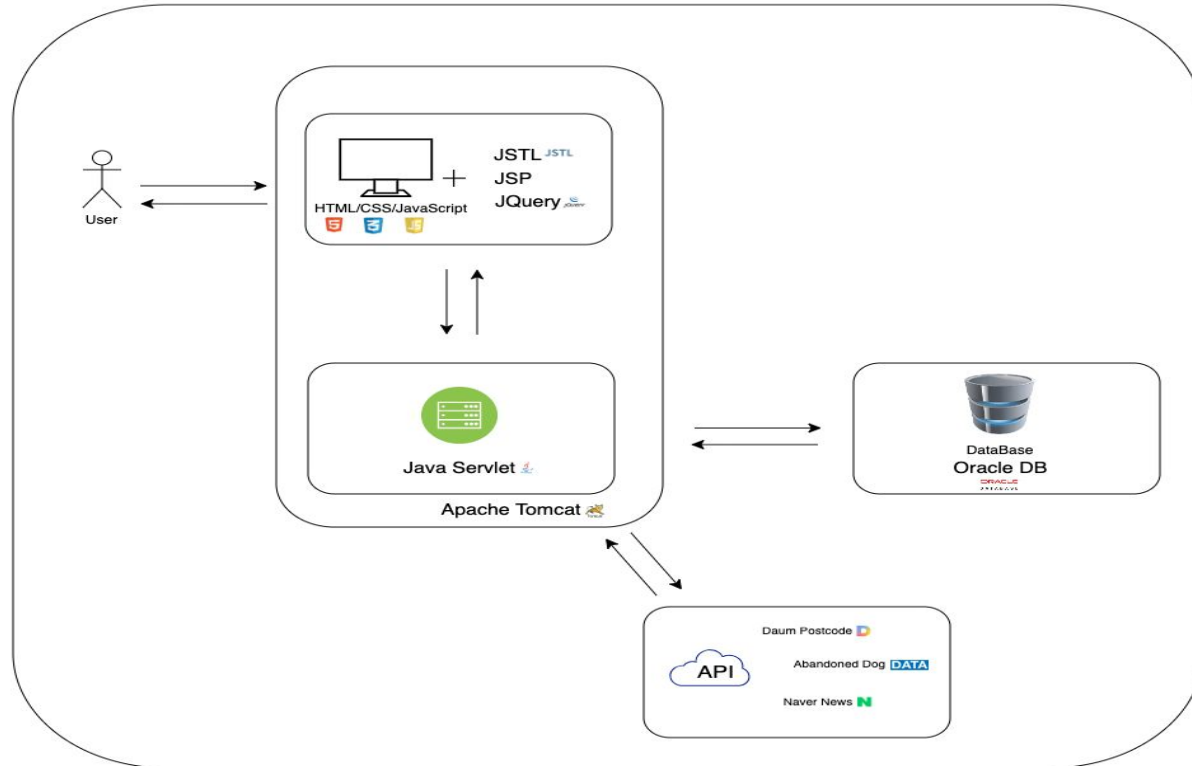
Development Environment

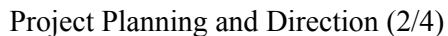
Languages	Java(JDK 17), HTML/CSS, JavaScript
Server	Apache Tomcat 9.0
Frameworks	Java Servlets & JSP
Database	Oracle DB
IDE	Eclipse, IntelliJ, SQL Developer
API, Libraries	Stray Dogs API, Daum Postcode API, Gson, JSoup, JSTL, JQuery



GitHub <https://github.com/h3lim/Pet-Community-Web>

System Architecture





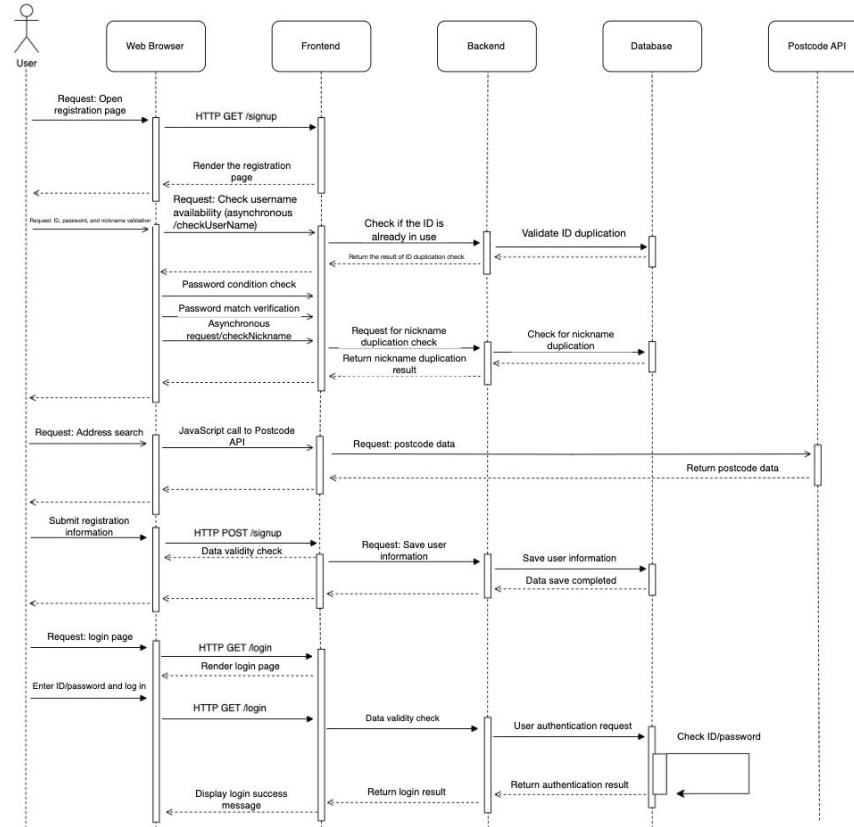
ERD Design

[illegible]



Project Planning and Direction (3/4)

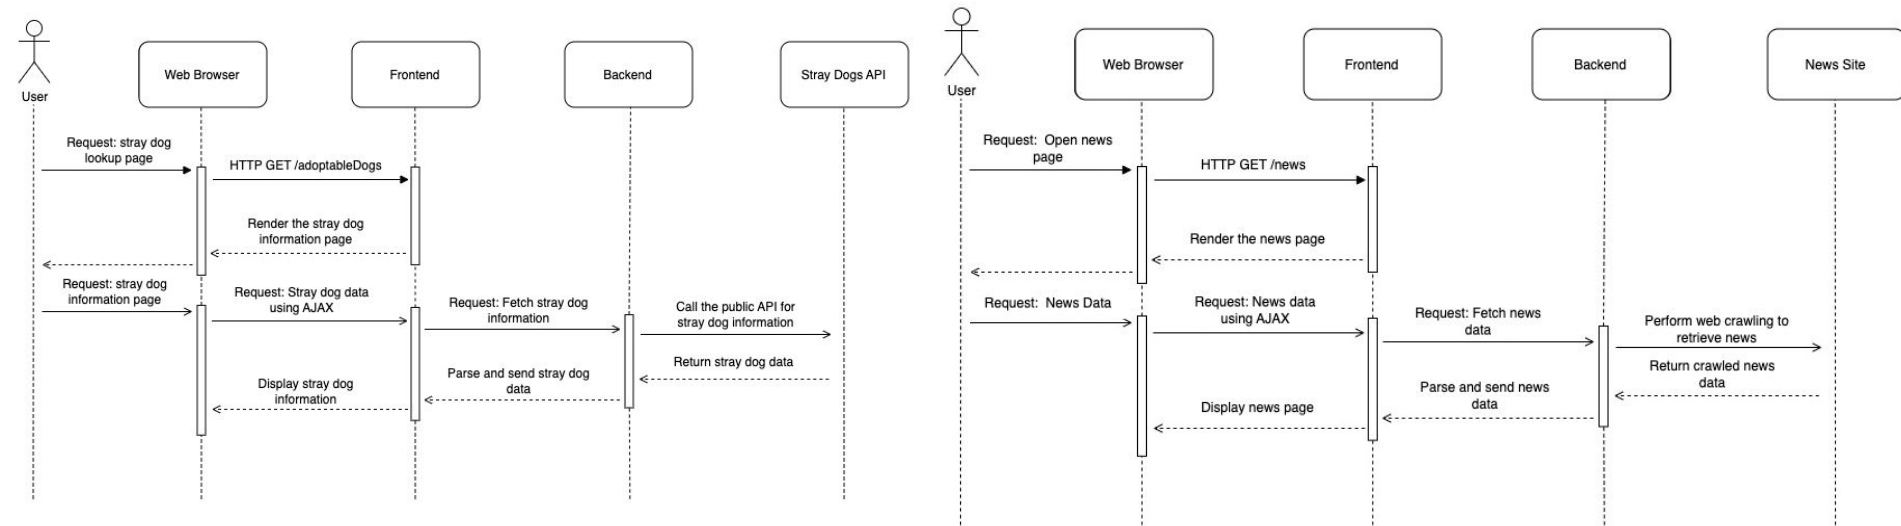
Sequence Diagram - User Registration and Login





Project Planning and Direction (3/4)

Sequence Diagram - Stray Dog Lookup, Pet-Related News Lookup





News Crawling Feature

- The system provides the latest pet-related news to users by crawling data from Naver News. To achieve this, the JSoup library is used to parse web pages and extract necessary information.

Implementation Details

- The NaverNewsCrawl class uses JSoup to fetch the HTML code of the Naver News page.
- Desired elements (e.g., article title, link, image, summary) are extracted using CSS selectors from the retrieved HTML document.
- Extracted data is stored in NewsVO objects and managed as a list.
- The CrawlServlet handles client requests for news and returns data in JSON format using Gson.
- The news.jsp page uses AJAX to fetch data from the servlet and dynamically display the news content.



Stray Dog Information Lookup Feature

- The system provides real-time stray animal information from the Public Data Portal to users. The DogDataFetcher class is used to communicate with the external public API to collect data and present it to users. The dog.jsp page uses AJAX to fetch data from the servlet and dynamically display the news content.

Implementation Details

- The DogDataFetcher class retrieves stray animal information from the Public Data Portal via HTTP requests.
- The service key and request parameters are included in the URL during the API call.
- The response XML data is converted to JSON, and JSONObject is used for parsing.
- Extracted data is stored in HashMap objects and managed as a list.
- The dog.jsp page allows users to input a start date and end date to query data and implements pagination for displaying the results.



API Rate Limit Issue with Naver News Crawling

When crawling data from Naver News search pages, there is a potential risk of hitting the server's maximum request rate, leading to rate-limit blocks. This is especially likely when multiple consecutive requests are made, which may cause the server to restrict or block the IP.

Symptoms

- After making several calls within a short time, data may fail to load, or crawling attempts may fail.
- Jsoup.connect() may return an IOException, or network block notifications may occur.

Cause Analysis

- Frequent requests may be interpreted as spam by the server.
- Insufficient delay between requests can result in exceeding the API rate limit.

Solutions

- Use `Thread.sleep(REQUEST_DELAY);` in the code to apply a 1-second delay between requests to mitigate rate-limit issues.
- Introduce caching to prevent redundant calls within a specific time frame, allowing frequently requested data to be reused and reducing API call frequency.



Response Delay and Intermittent Errors in Stray Dog API

Delays or intermittent errors may occur when calling the public API, resulting in failure to retrieve data.

Symptoms

- Specific error messages frequently appeared on the page during intermittent errors, leading to repeated error displays if unresolved.
- Long response times from the API resulted in prolonged page loading times, with some requests eventually failing due to timeouts.

Cause Analysis

- Public APIs may have variable server status and response times based on usage, with potential abnormal states.

Solutions

- Implement retry logic with specified MAX_RETRIES and RETRY_DELAY_MS values to perform re-requests after a delay in case of a failure.
- Limit retry attempts to 3 to avoid excessive requests and set a delay time of 2 seconds to prevent server overload.